



local-ai-lab

Build local AI, one lesson at a time.

RAG · MCP · LANGCHAIN · LANGGRAPH

AI Developer Cheat-Sheet

Light system. Heavy thinking. A one-stop quick reference for building with AI — locally and with cloud agents. Prepare your machine, pick the right model for the job, validate your work, and ship. Part of [local-ai-lab](#) by [Nik Reljin](#).

Model names move fast. Cloud/media/local model names below are current **as of June 2026** and every row links to the source — always confirm the exact version at the link. Claude model IDs are exact and stable.

1 · Prepare your local (Ollama + BitNet.cpp)

Run models on your own hardware — no API keys, no data leaving the box. Ollama for general local LLMs; BitNet.cpp for ultra-light 1-bit models that run on CPU / older machines.

Ollama — [ollama.com](#)

Step	Command
Install (Linux)	<code>curl -fsSL https://ollama.com/install.sh sh</code>
Start the server	<code>ollama serve</code> (API on <code>http://localhost:11434</code>)
Pull a model	<code>ollama pull qwen2.5-coder:7b</code>
Chat in terminal	<code>ollama run llama3.3</code>
List / remove	<code>ollama list · ollama rm <model></code>
Use the API	<code>curl http://localhost:11434/api/generate -d '{"model":"llama3.3","prompt":"hi"}'</code>
GPU / tuning	<code>OLLAMA_NUM_GPU, OLLAMA_HOST, OLLAMA_KEEP_ALIVE</code> env vars

Pick-and-run a model from a simple UI with [ai-runner](#).

BitNet.cpp — [github.com/microsoft/BitNet](#)

1-bit (ternary -1/0/+1) inference, ~0.4 GB RAM for a 2B model, CPU-friendly. Based on `llama.cpp`.

Step	Command
Clone + submodules	<code>git clone --recursive https://github.com/microsoft/BitNet && cd BitNet</code>
Build	<code>pip install -r requirements.txt && python setup_env.py -md models -q i2_s</code>
Get a model	<code>huggingface-cli download microsoft/bitnet-b1.58-2B-4T-gguf --local-dir models/BitNet-b1.58-2B-4T</code>
Run inference	<code>python run_inference.py -m models/.../ggml-model-i2_s.gguf -p "Hello" -cnv</code>

Producing your own small models? See [shrink-llm](#) (quantize/prune/distill) and [finetorch](#) (LoRA/QLoRA).

2 · Coding agents · testing · validating PRs

Pick a coding agent

Agent	Best at	Local models?	Notes
Claude Code	Long-horizon agentic coding, refactors, reviews	Cloud (Claude); local via proxy	Default <code>claude-opus-4-8</code> ; <code>/code-review</code> , <code>/security-review</code> built in
OpenAI Codex (CLI)	Agentic "Goal Mode", PR creation, frontend	Cloud (GPT-5.x)	Runs on GPT-5.5 / GPT-5.2-Codex; MCP support
Gemini CLI	Large-context, vision-assisted, free tier	Cloud (Gemini 2.5)	Strong on huge context + image input
Ollama + your editor	Fully offline, private	Yes — <code>qwen2.5-coder</code> , <code>llama3.3</code>	No keys, no egress; pair with Continue/aider

Run the tests (this repo's pattern)

```
python -m venv venv && source venv/bin/activate
pip install -r requirements.txt
pytest -q                # offline tests – no network, no LLM
./run -l N test          # run a lesson's tests via the harness
```

- Keep unit tests **offline** (no network, no live model) so they're fast and deterministic.
- Let an agent **write and run** the tests, then read the failures — don't trust "looks right".

PR validation checklist

- [] Tests green locally (pytest -q) and in CI
- [] **Conventional Commits** — feat: fix: chore: docs: refactor: ci:
- [] Ran /code-review (bugs) and /security-review (secrets, injection) on the diff
- [] No secrets committed — .env ignored, placeholders in .env.example
- [] CHANGELOG updated if the release workflow keys off it
- [] **No AI-attribution trailers** — author is you; no "Co-Authored-By: ...AI", no "Generated with..."

3 · Which model for what (the picker)

Match the **task** to a **cloud** pick and a **local** equivalent. Cloud = best quality / least setup; local = private / free / offline.

Text — code, chat, reasoning

Task	Cloud (service · model)	Local (Ollama / BitNet)
Agentic coding	Claude claude-opus-4-8 · docs	qwen2.5-coder:32b
Fast / cheap code & chat	Claude claude-sonnet-4-6 / claude-haiku-4-5; Gemini 2.5 Flash	qwen2.5-coder:7b, llama3.3:8b
Hardest reasoning	Claude claude-fable-5; OpenAI GPT-5.5; Gemini 2.5 Pro	llama3.3:70b
CPU-only / tiny footprint	—	BitNet bitnet-b1.58-2B-4T

Vision — analyze / describe images

Task	Cloud	Local
Read content in an image, charts, docs	Gemini 2.5 Pro/Flash (vision); Claude vision (claude-opus-4-8)	qwen2.5-vl, llama3.2-vision, gemma3
OCR / document extraction	Gemini 2.5 Pro; Claude (PDF input)	qwen2.5-vl + Tesseract

Embeddings — for RAG / search

Task	Cloud	Local
------	-------	-------

Embed text for retrieval	OpenAI <code>text-embedding-3</code> ; Gemini embeddings	<code>nomic-embed-text</code> (default), <code>bge-m3</code>
--------------------------	---	---

Media generation (links — verify versions)

Make...	Service · model	Link
Images	Gemini 2.5 Flash Image ("Nano Banana") · FLUX.2 · GPT-Image · Imagen 4	nano-banana · FLUX
Images (local)	SDXL / FLUX via ComfyUI / Diffusers	ComfyUI
Video	Veo 3.1 (audio) · Kling 3.0 (cheap) · Runway Gen-4.5 (control) · Seedance 2.0	Veo · Runway
Speech → text (STT)	OpenAI Whisper API · Gemini	Whisper (local), <code>faster-whisper</code>
Text → speech (TTS)	ElevenLabs · Gemini TTS	Kokoro , Piper (local)

Default to the **latest Claude** model for serious code/agent work; reach for Gemini for big-context + vision, and local models when privacy/offline/cost matters.

4 · RAG — do / don't

Pipeline (see `localrag/`): **extract** → **chunk** → **store/embed** → **retrieve** → **generate (grounded)**.

Do	Don't
Right-size chunks; keep overlap small and consistent	Dump the whole corpus into the prompt
Retrieve top-k, then cite sources in the answer	Answer ungrounded / without showing sources
Use hybrid (BM25 + embeddings) for recall (Lesson 3)	Rely on a single retrieval signal
Evaluate retrieval + answers, regression-test (Lesson 5)	Ship without an eval / "it looked fine"
Treat retrieved text as data , sanitize it (Lesson 4)	Treat retrieved text as instructions (injection)
Prompt-cache large stable context to cut cost/latency	Re-send the same big context uncached every call

5 · MCP & tool design — do / don't

Expose your retriever/tools to an LLM via **MCP** (see `mcp_server.py` — FastMCP over stdio). Claude can also connect to MCP servers via the `mcp_servers` API param.

Do	Don't
Clear names; prescriptive descriptions ("call this when...")	Vague names / what-it-does-only descriptions
Typed JSON schema; <code>enum</code> for fixed sets; mark required	Untyped free-form string inputs
Return errors with <code>is_error: true</code> + a useful message	Swallow errors or return silent failures
Keep the tool set small and focused	Expose dozens of overlapping tools
Gate side-effecting / irreversible tools	Auto-run destructive actions unprompted
Keep secrets server-side; pass via env/vault	Put API keys in prompts, args, or tool output

6 · LangChain / LangGraph / Lang* — when to use what

Tool	Use it for	Don't
LangChain	Glue: chains, loaders, retrievers, model adapters	Wrapping a single one-shot call in 5 abstractions
LangGraph	Stateful agents / multi-step graphs, retries, branching	Linear pipelines a plain function handles
LangSmith	Tracing, eval, debugging prod chains	As a hard dependency for tiny scripts
LlamaIndex	Heavy RAG indexing / query engines	When <code>localrag</code> + a vector store already fits

All of the above point at **local Ollama** via the OpenAI-compatible endpoint (`http://localhost:11434/v1`) — develop offline, swap to cloud later. **Don't over-abstract:** if it's one model call, just call the model.

7 · Shrinking LLMs for local use

Make big models fit small machines. See [shrink-llm](#).

Technique	What it does	Trade-off
Quantization (GGUF q4/q5/q8)	Fewer bits per weight	q4 $\approx$ 4× smaller, slight quality drop; q8 near-lossless
Pruning	Drop low-impact weights	Needs fine-tune to recover quality

Distillation	Train a small "student" from a big "teacher"	Up-front training cost
1-bit (BitNet)	Native ternary weights	Smallest/fastest on CPU; use the native b1.58 models

Rule of thumb: start with a q4/q5 **GGUF** of a good model in Ollama; drop to **BitNet** when you need CPU-only or sub-GB footprint; reach for pruning/distillation/LoRA only when a stock model won't fit or won't perform.

Links

- **local-ai-lab** course & code — <https://github.com/nikolareljin/local-ai-lab>
- Ollama — <https://ollama.com> · Models — <https://ollama.com/library>
- BitNet — <https://github.com/microsoft/BitNet>
- Model Context Protocol — <https://modelcontextprotocol.io>
- Claude Platform docs — <https://platform.claude.com/docs>
- Google AI (Gemini) — <https://ai.google.dev>
- shrink-llm — <https://github.com/nikolareljin/shrink-llm> · finetorch — <https://github.com/nikolareljin/finetorch> · ai-runner — <https://github.com/nikolareljin/ai-runner>

local-ai-lab — built and maintained by **Nik Reljin**. Cheat-sheet current as of June 2026; verify fast-moving model versions at the linked sources.