

Lesson 3 · Hybrid Retrieval & Reranking

Install (Linux · macOS · Windows): [guide](#)

Part of [local-ai-lab](#) — a hands-on course for building local AI.

Course home: <https://nikolareljin.github.io/local-ai-lab/> **Source:** <https://github.com/nikolareljin/local-ai-lab> **Author:** [Nik Reljin](#)

Lessons: 1 · [RAG](#) → 2 · [MCP](#) → 3 · **Hybrid retrieval (you are here)** → 4 · RAG safety → 5 · RAG evaluation → 6 · Repo assistant → 7 · LangChain → ... → 15 · Docs from changes

Status: working demo. Runnable in **Python, Node.js, and C# / .NET** — same algorithm, three languages, identical rankings. Runs 100% offline, no embedding model required. See **From demo to production** at the end for what to harden for real use.

What you'll learn

[Lesson 1](#) put **BM25** and **embeddings** side by side as two ways to retrieve. This lesson answers the question that always comes next: **which one should I use?** The honest answer is **both** — and this lesson shows why, and how to combine them.



By the end you'll understand:

- **BM25** — best for exact terms: IDs, error codes, names, keywords
- **Embeddings** — best for semantic similarity and paraphrasing
- **RRF** (Reciprocal Rank Fusion) — combine multiple ranked lists into one
- **Reranking** — reorder the top chunks before generation for precision
- **Query rewriting** — turn a vague question into a better retrieval query
- **Metadata filtering** — narrow by file type, product, date, page, module, system

The one idea: BM25 and embeddings are not enemies. They solve **different** retrieval problems. "**What fixes error E_4096?**" needs exact-term BM25; "**why won't it turn on?**" needs semantic matching. Good enterprise RAG runs both and fuses the results.

The demo

A tiny corpus of three support docs lives in `data/`. We run **two queries** that pull in opposite directions:

- an **exact-keyword** query — `error E_4096`

- a **keyword-free paraphrase** — `broken gadget` (no document contains either word)

...and print the BM25 ranking, the semantic ranking, and the fused (RRF) ranking for each.

Run the comparison (offline, no dependencies)

From the repo root — the `demo` action prints the rankings and exits, with **nothing to install** (pure standard library). Pick any language; all three give the **same** output:

```
./run -l 3 demo          # Python — print the BM25 / semantic / hybrid comparison and exit
./run -l 3 --lang node demo # Node.js — same output
./run -l 3 --lang csharp demo # C# / .NET 8 — same output
./run -l 3 test          # the offline Python test
./run -l 3 show          # walk through this lesson's steps (code, data, prompts, commands)
```

`./run -l 3 demo` is the dependency-free **comparison**: it prints the rankings for two queries and exits — no server, no install, byte-identical across Python, Node.js and C#.

Experiment in the GUI (needs Flask)

For a hands-on feel, bare `./run -l 3` opens an interactive **experiment GUI** over the bundled short story — **The Voyage of Caretta the Magnificent** (the magic turtle who became an astronaut), five chapters in `story/`. Type a query and watch the BM25, semantic and hybrid rankings (and the numbers behind them) recompute as you tune the knobs — no code editing.

```
./run -l 3          # opens http://127.0.0.1:<port> — the experiment GUI over the 5-chapter story
```

The GUI is a small Flask app, so unlike the `demo` it needs one dependency. `./run` installs it into the project venv automatically on first use (`pip install -r requirements.txt` if you prefer manual).

Try these and watch the two retrievers diverge:

Query	Type	Expect
Nuevo Edén	exact name	BM25 nails the discovery chapter (04-nuevo-eden.md)
Alpha Centauri	exact name	the chapters that name the star system rank top
who found the new planet	paraphrase	the semantic arm surfaces the discovery chapter with no shared keyword
how did they keep the turtle alive in space	paraphrase	the spacesuit chapter (02-spacesuit.md) rises

Exact names are BM25's strength; the paraphrases are where the semantic arm earns its keep — and the hybrid (RRF) column gets both right. Prose like this shows the trade-off far better than the toy support-doc corpus.

Output:

```
Query: "error E_4096"
BM25 (lexical): ['error_codes.md']
Semantic (stand): ['error_codes.md']
Hybrid (RRF): ['error_codes.md']

Query: "broken gadget"
BM25 (lexical): []
Semantic (stand): ['power_issues.md']
Hybrid (RRF): ['power_issues.md']
```

Notice the second query: **BM25 returns nothing** — no document literally contains **broken** or **gadget** — yet the semantic arm recovers `power_issues.md`, and the hybrid keeps that hit. That is the divergence fusion exists for.

Why no embedding model? To keep the demo offline and dependency-free, the "semantic" arm is a small **synonym-expanded overlap** stand-in — enough to behave like embeddings (mapping **broken** → **fail/dead**, so it recovers the **power/startup** doc that BM25 misses entirely). In production you swap in real sentence embeddings; the **fusion** code doesn't change. That swap is the first item in **From demo to production**.

Concept 1 · BM25 — lexical retrieval

BM25 ranks documents by **exact term overlap**, weighting rare terms more (IDF) and damping repeated terms and long documents. The compact implementation in the demo:

```
idf = math.log(1 + (n - df[t] + 0.5) / (df[t] + 0.5))      # rare terms count more
tf = d["tokens"].count(t)
score += idf * (tf * (K1 + 1)) / (tf + K1 * (1 - B + B * dl / avgdl))  # tf, length-normalized
```

- **Wins** the exact query: `E_4096` appears in exactly one document, so its IDF is high and `error_codes.md` shoots to the top.
- **Losses** on a keyword-free paraphrase: ask **"broken gadget"** and BM25 returns **nothing** — no document shares a single token with the query.

Teaching point: reach for BM25 when the user knows the exact string — an error code, a product SKU, a function name, a person's name. It is fast, needs no model, and is unbeatable at literal matching.

Concept 2 · Embeddings — semantic retrieval

Embeddings map text into a vector space where **meaning** is distance, so paraphrases land near each other even with no shared words. In the demo a synonym-expanded overlap **stands in** for this:

```
SYNONYMS = {"broken": ["fail", "fails", "dead"], "turn": ["power", "start", "startup", "boot"], ...}
# "broken gadget" now overlaps the power/startup document (broken -> fail/dead)
```

- **Wins** the keyword-free paraphrase: `power_issues.md` surfaces even though "**broken gadget**" shares no word with any document.
- **Loses** on exact IDs: a rare code like `E_4096` can get diluted among "semantically similar" text.

Teaching point: reach for embeddings when the user describes a problem in their own words. They cost more (you embed every chunk and the query) but they recover the recall BM25 misses.

Concept 3 · RRF — fusing the two rankings

Each retriever produces a different **ranked list**. **Reciprocal Rank Fusion** combines them without needing comparable scores — each list contributes $1 / (k + \text{rank})$ to a document, then you re-sort:

```
fused[name] += 1.0 / (RRF_K + position + 1)    # RRF_K = 60
```

Because it works on **ranks, not raw scores**, BM25 and embeddings get an equal vote despite living on totally different scales. In the demo, the fused ranking gets **both** queries right where each single retriever gets only one right.

Teaching point: RRF is the cheapest, most robust way to combine retrievers — a few lines, no tuning of score scales. It's the default first step of any hybrid system.

Concept 4 · Reranking, query rewriting & metadata filtering

Three more levers that production systems add on top of fusion:

- **Reranking** — take the fused top-k and reorder them with a stronger model (a cross-encoder that scores each `(query, chunk)` pair). Retrieval optimizes **recall**; reranking optimizes **precision** at the very top, where the LLM actually reads.
- **Query rewriting** — rewrite a vague question into a better retrieval query **before** searching ("**it keeps crashing**" → "**device repeatedly powers off overheating E_4096**").
- **Metadata filtering** — restrict candidates by structured fields (file type, product, version, date, page, module) before or after ranking, to cut noise in large mixed corpora.

Teaching point: reranking is usually the single biggest quality win — and the main added cost and latency. Add it once you've proven fusion isn't enough (and measure it — see Lesson 5 · RAG evaluation).

When each technique wins

Technique	Best for	Weak at	Cost	Needs a model?
-----------	----------	---------	------	----------------

BM25	exact terms: IDs, error codes, names	paraphrase, synonyms	very low	no
Embeddings	semantic similarity, paraphrasing	exact tokens, rare IDs	medium	yes
RRF (fusion)	combining BM25 + embeddings	nothing — it's combinatorial	negligible	no
Reranking	precision at the top-k	adds latency; needs top-k first	higher	yes
Query rewriting	vague / underspecified questions	can over-expand	one LLM call	yes
Metadata filtering	scoping large mixed corpora	needs good metadata	negligible	no

On the demo corpus BM25 nails the exact code but returns **nothing** for the keyword-free paraphrase, while the semantic arm recovers the right document — and RRF keeps whichever arm found the answer. That's the whole argument for fusion: you never lose one retriever's strength.

Polyglot by design

The retrieval algorithm is language-agnostic, so this lesson ships in **Python, Node.js, and C# / .NET** — each dependency-free, each reading the same `data/`, each producing byte-identical rankings. Compare the three implementations side by side: the **concepts** (BM25, synonym expansion, RRF) are the same; only the syntax changes.

Port	Entry point	Run
Python	python/hybrid_demo.py	<code>./run -l 3 demo</code> · <code>./run -l 3 test</code>
Node.js	node/hybrid_demo.mjs	<code>./run -l 3 --lang node demo</code>
.NET 8	dotnet/Program.cs	<code>./run -l 3 --lang csharp demo</code>

All three commands are declared once in `lesson.json` — the single source of truth the `./run` engine reads. `./run -l 3 show` renders this lesson's elements (notes, code, sample data, prompts, commands) in order.

Exercises

- **Real embeddings:** replace the semantic stand-in with a local sentence-embedding model and rerun both queries — does fusion still help?
- **Tune RRF:** change `RRF_K` and the per-arm weighting; observe the effect on the fused order.
- **Add a third retriever:** fold in a title/metadata match as a third ranked list — RRF takes any number of lists.
- **Break it:** add a document that mentions **E_4096** in passing and watch how BM25 vs hybrid handle the near-duplicate.

From demo to production

- **Real embeddings** for the semantic arm (a local sentence-embedding model, or the repo's optional embedding path) — the demo's synonym stand-in is only to stay offline.
- **A cross-encoder reranker** over the fused top-k — usually the biggest single quality win, at some latency/compute cost.
- **Query rewriting / expansion** as an LLM step before retrieval for vague queries.
- **A metadata-rich index** so you can filter by product, version, date, and module.
- **Tune `RRF_K`, the top-k, and per-arm weights** per corpus.
- **Measure it** — pair with **Lesson 5 · RAG evaluation** (coming) and prove hybrid actually beats single retrievers on your golden set before shipping it.

Next lesson

Lesson 4 · RAG safety & prompt injection — retrieved documents are untrusted input; never follow instructions found inside them.

Course: nikolareljin.github.io/local-ai-lab · Author: **Nik Reljin**