

Lesson 8 · A Stateful Agent with LangGraph

PDF: [this lesson](#) · Install (Linux · macOS · Windows): [guide](#) · [PDF](#)

Part of [local-ai-lab](#) - a hands-on course for building local AI.

Course home: <https://nikolareljin.github.io/local-ai-lab/> Source: <https://github.com/nikolareljin/local-ai-lab>

Lessons: 1 · [RAG](#) → 2 · [MCP](#) → 3 · [Hybrid retrieval](#) → 4 · [RAG safety](#) → 5 · [RAG evaluation](#) → 6 · [Repo assistant](#) → 7 · [LangChain](#) → 8 · **LangGraph (you are here)** → 9 · [Ollama tools](#) → 10 · [Semantic Kernel](#) → 11 · [Bedrock Agents](#) → 12 · [Google ADK](#) → ... → 15 · [Docs from changes](#)

Status: **planned**. Outline below; full step-by-step coming later. Star the repo to follow along.


The idea

A straight RAG chain (Lessons 1 and 7) is **linear**: retrieve → ground → answer, once. Real assistants need **control flow** - decide whether to retrieve at all, retry with a reformulated query when results are weak, call more than one tool, and remember the conversation. **LangGraph** models this as a **stateful graph** of nodes and edges with explicit state.

This lesson turns the RAG pipeline into an agent that **thinks about its own retrieval**.

What you'll build

A graph with nodes such as:



What you'll learn

- **State** - a typed graph state that flows between nodes (question, docs, draft, attempts)
- **Nodes & edges** - retrieval, grading, query rewriting, and generation as composable steps

- **Conditional edges** - branch on "are these results good enough?" and loop to self-correct
- **Cycles** - corrective RAG: re-retrieve with a better query instead of answering badly
- **Memory & checkpoints** - persist state so a conversation survives across turns
- **Tool routing** - combine the Lesson 2 MCP `search_docs` tool with other tools

Builds on the whole course

Concept	Comes from
retrieve / ground / answer	Lesson 1
<code>search_docs</code> as a callable tool	Lesson 2
LangChain retrievers & chat models	Lesson 7
graph, state, cycles, self-correction	Lesson 8 (this one)

Where it all lands: by the end of the course you can build a local, private, self-correcting document agent - and explain every layer, from `pypdf` extraction up to a LangGraph control loop.

Prerequisites

Lessons [1](#), [2](#), and [7](#) · [LangChain](#). The agent reuses the retriever, the MCP tool, and the LangChain components from earlier lessons.

Next lesson

Lesson 9 · Ollama + Function Calling → - give a local model its own tools, fully offline, then carry the same `search_docs` capability through Semantic Kernel, Bedrock, and ADK.

Course: nikolareljin.github.io/local-ai-lab · **Author:** [Nik Reljin](#)